



Xi'an Jiaotong-Liverpool University

西交利物浦大學

# Welcome to Microprocessor Systems


An introduction to CPT210

# Module Information

About the module instructor, teaching organisation.

# Teaching Organisation

- Module instructor: Jianjun Chen
  - Email address: [Jianjun.Chen@xjtlu.edu.cn](mailto:Jianjun.Chen@xjtlu.edu.cn)
  - Office: SD 541
  - Office hours: please check the module page online
    - If you wish to see me at other times, please email me first.
- Module instructor: Kok Hoe Wong
  - Email address: [kh.wong@xjtlu.edu.cn](mailto:kh.wong@xjtlu.edu.cn)
  - Office: SD 431
  - Office hours: please check the module page online

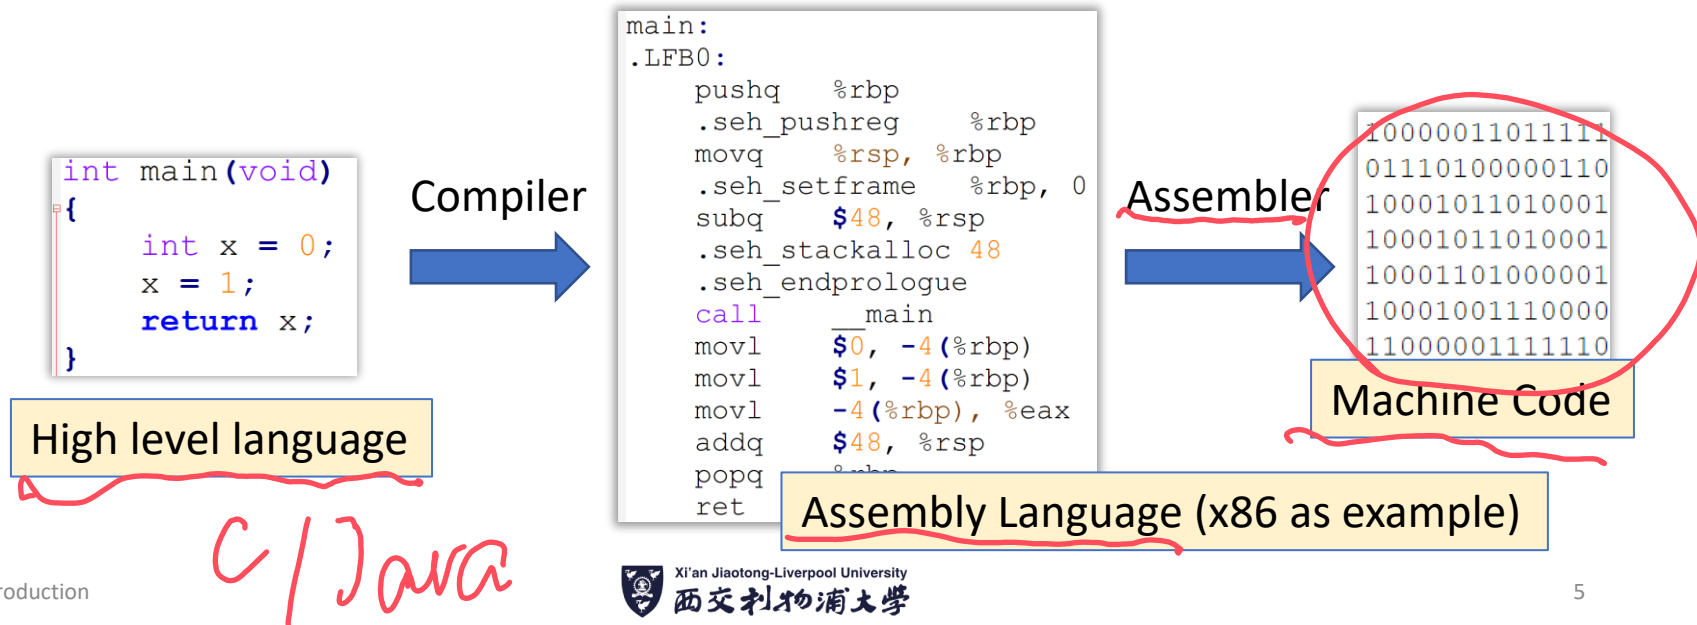


# Teaching Organisation

- Support:
  - I have office hours every week.
  - Work in groups with your friends during labs. Discussions with classmates enable you to better understand concepts and terminologies.
  - I will add one or more notice boards on the Learning Mall about issues (like coursework, exam, labs) related to this module, I will also send notifications to you about any updates. **Please check emails frequently.**
- Grades: Please check the information about coursework and exams on e-bridge.
  - You are encouraged to discuss ideas (Not solutions!) with others when doing coursework.
  - But your assignment submissions must be your own works.

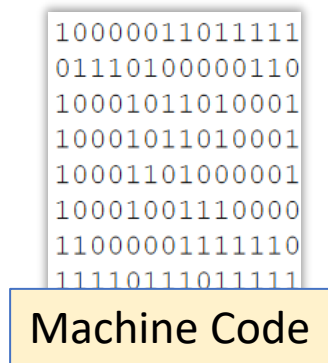
# About this Module

- This module is a “bridge” between digital electronics (hardware) and high level languages (software).
- You will learn basic concepts required to understand how computers work
  - What programs are? How does your C/C++ code become the instructions that your computer understands?
  - How these instructions are executed in hardware?



# About this Module

- This module is a “bridge” between digital electronics (hardware) and high level languages (software).
- You will learn basic concepts required to understand how computers work
  - What programs are? How does your C/C++ code become the instructions that your computer understands?
  - How are these instructions executed in hardware?



# Focus of this Module

- We will focus on Instruction Set Architectures and their use through **Assembly Language Programming**.
- The architecture studied in this module will be **ARM**.

# Introduction to Microprocessor Systems

Evolution of technology

Design of ISA





# The Evolution of Computers

- Electronic Delay Storage Automatic Computer ([EDSAC](#))
  - First [practical](#) programmable computer, by UK. In Cambridge 1949.
  - On average, processes 650 instructions per second.
  - 1024 17-bit words of memory in [mercury ultrasonic delay lines](#).
  - 3000 valves, 12 kW power consumption, occupied a room 5m by 4m.
  - Early use to solve problems in meteorology, genetics and X-ray crystallography.



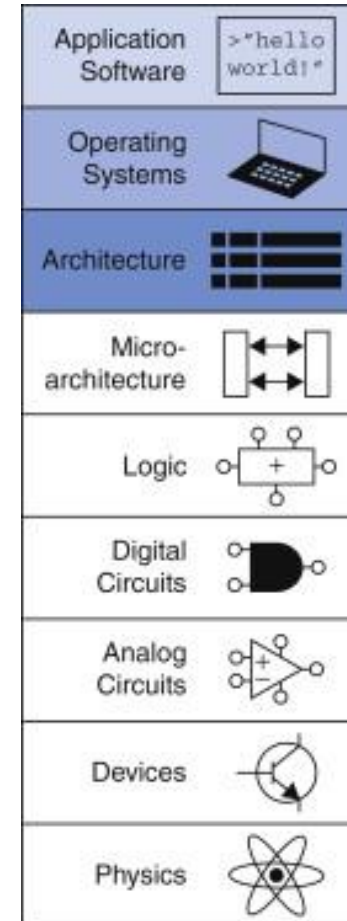
# The Evolution of Computers

- ARM7
  - Up to 130 million instructions per second. 1995-2011
  - One of the most successful embedded processor.
- The picture shows LPC-H2124 header board. 
  - 256KB of program flash
  - 16KB of RAM.
  - The square chip in the middle is the microcontroller
- Original ARM design:
  - Steve Furber, Acro [RISC](#) Machine
  - Cambridge, 1985
- Evolution of computer architecture



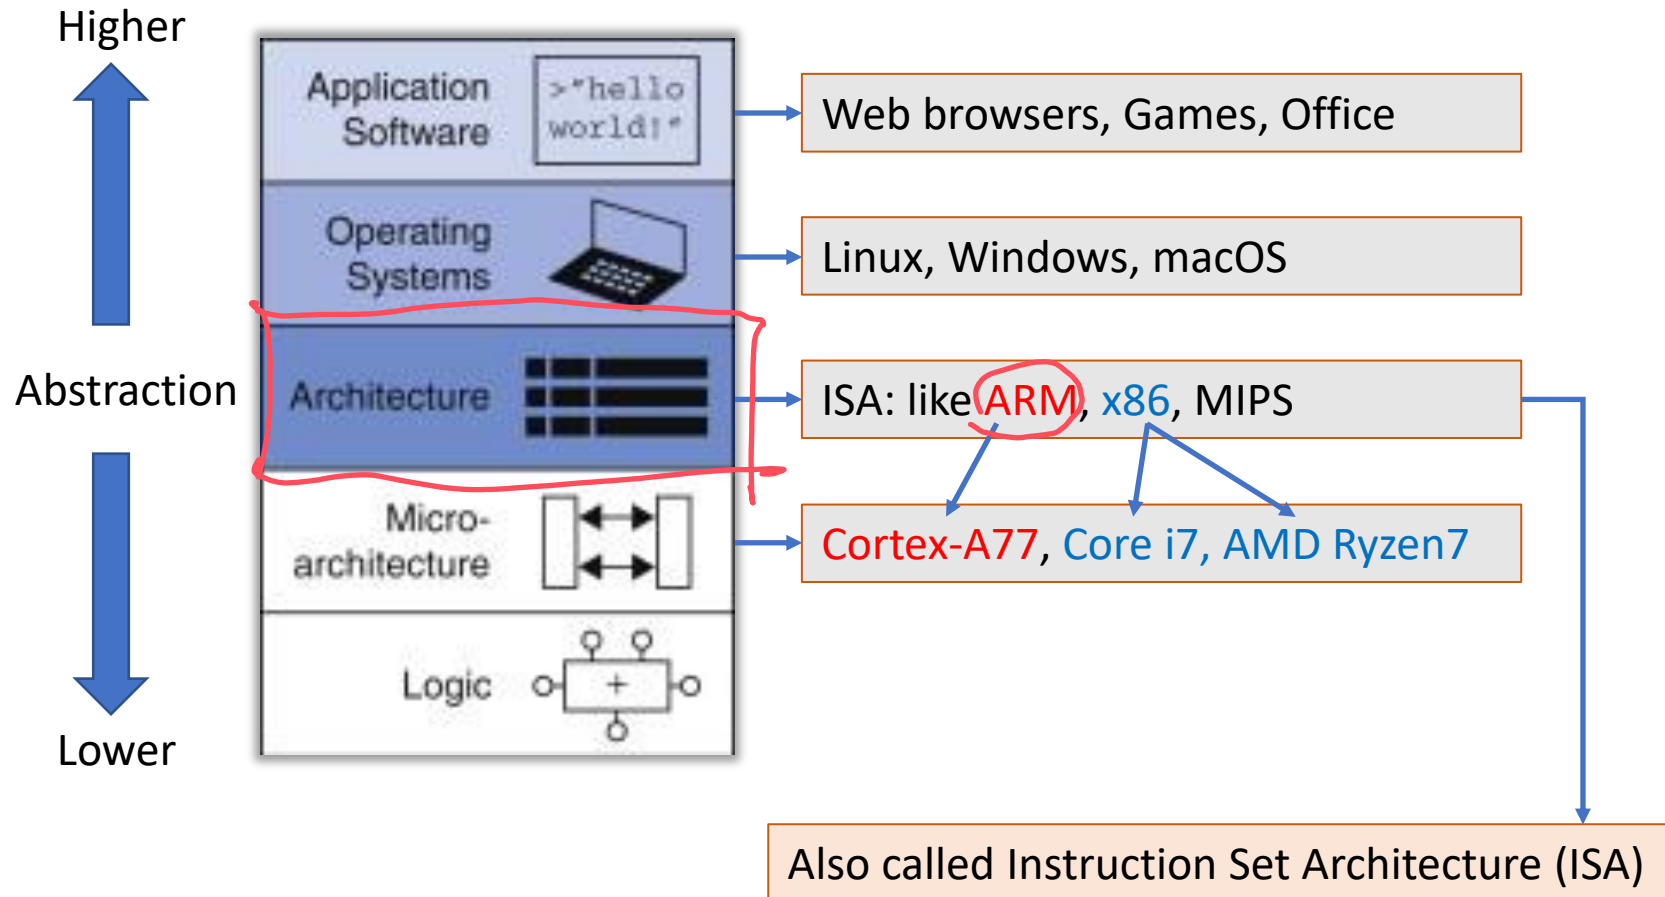
# What is Computer Architecture?

- *Computer architectures* represent the means of interconnectivity for a computer's hardware components as well as the mode of data transfer and processing exhibited.
  - Paul J. Fortier, Howard E. Michel, in [\*Computer Systems Performance Evaluation and Prediction\*](#), 2003
- The *architecture* is the programmer's view of a computer. It is defined by the instruction set (language) and operand locations (registers and memory). Many different architectures exist, such as ARM, x86, MIPS, SPARC, and PowerPC.
  - Sarah L. Harris, David Money Harris, in [\*Digital Design and Computer Architecture\*](#), 2016



# What is Computer Architecture?

Architecture is about design, not just hardware.



# Instruction Set Architecture (ISA)

- Instruction (or Operation Code) set.
  - ADD, SUB, RSB, etc.
- Organisation of programmable storage.
  - Where data is stored.
- Mode of addressing and accessing data items and instructions.
  - How data is stored/retrieved.
- Behaviour on exceptional conditions, e.g.:
  - Hardware divide by zero.
  - An interrupt occurred
- ISA is about what the computer does (not how it does).

# Instruction Set Architecture (ISA)

- Examples of ISAs:
  - **8086/Pentium (x86) ISA:** Widely used, compatible with a lot of operating systems and software.
  - **ARM ISA:** Supports highly optimising compiler & operating system software for embedded applications.

## x86 instruction

<b>ADD</b>	Add numbers
<b>AND</b>	Logical AND
<b>MOV</b>	Copy data between memory blocks/registers

## ARM instruction

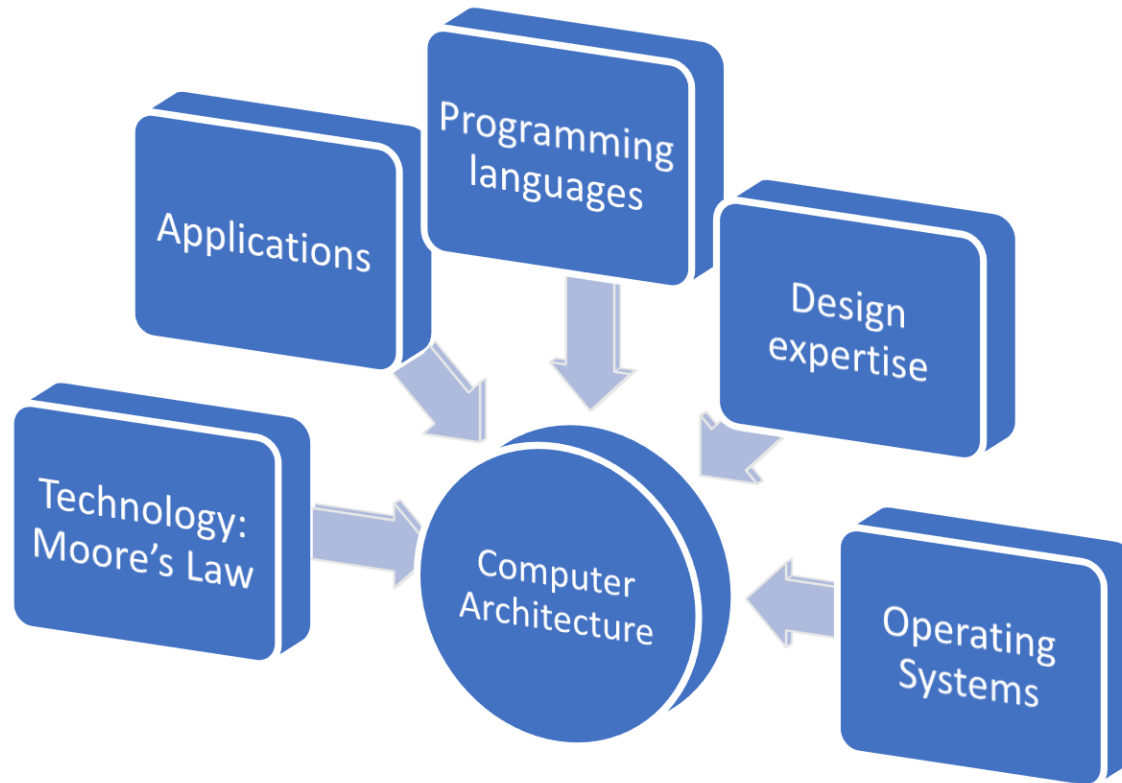
<b>ADD</b>	Add numbers
<b>AND</b>	Logical AND
<b>LDR</b>	Load from memory (can't save to memory)
<b>STR</b>	Store to memory (can't load from memory)

- Different implementations of the same ISA can run identical software.
  - Your compiled C program can run on any x86 computer as long as the operating system is the same and libraries are satisfied.

# Discussion

- Discuss the questions below based on your understanding and knowledge so far.
- Computer architecture affects:
  1. CPU design?
  2. Compiler design?
  3. Energy efficiency of the computer?
  4. Maximum memory size?
  5. UI design of the operating system (OS)?
  6. Security of the OS?

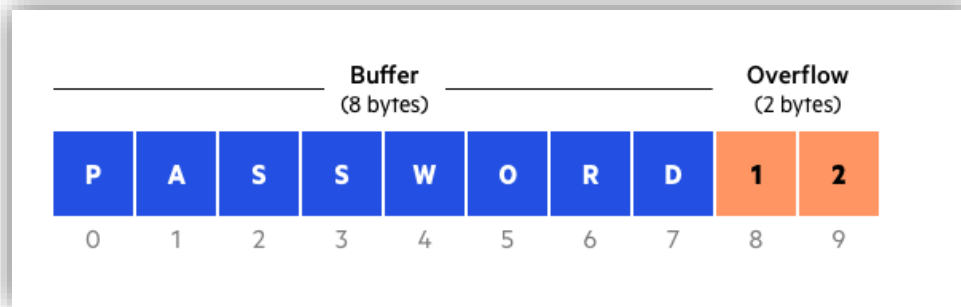
# What Factors Influence Computer Architecture?





# Example: Buffer Overflow Attack


- This attack targets programs that accept input data from untrusted sources and do not verify the length of the stored data.



- Overloaded buffer can override executable code with malicious code.
- Execution Disable Bit (EDB), introduced by Intel
  - Classifies areas in memory where a code can execute or not execute.
  - Stops an attack from running code in a non-executable region.
  - Does not shutdown the buffer overflow attack completely, but can help reduce the impact.


# Example: Application

- Special units for AI learning in CPU.
  - Implemented in recent chips like Snapdragon from Qualcomm, Apple Silicon, Kirin chips, Exynos chips.
- Video processing unit.
  - Intel chips provide video encoding/decoding accelerating
  - Can greatly speed up video decoding and encoding.



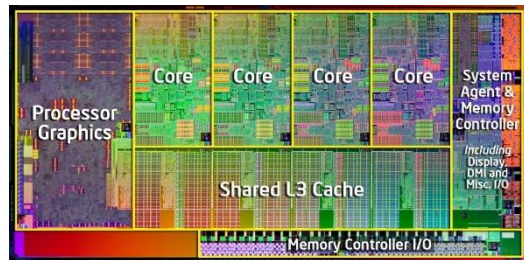
# Example: Technology (Moore's Law)

- In 1965, Gordon E. Moore, co-founder of Intel, postulated that the number of transistors that can be packed into a given unit of space will double about every two years. (See on [Wikipedia](https://en.wikipedia.org/wiki/Moore's_Law))





# Example: Technology (Moore's Law)

- VLSI (very large scale integration) has increased speed and density of CPUs by shrinking dimensions
  - This is limited by size of atoms (Quantum effects), so will stop.
- Moore's Law predicted a doubling of computing power every two years.
  - This has held true but will stop soon without a change in technology or architecture (Quantum computing).



i7-2600k: 4 cores

i7-12700k:  
8 P cores  
8 E cores



# Micro-architectures of ARM

- There are many architectures developed by ARM.
- Each architecture has many implementations (micro-architectures).
- Each micro-architecture has its own targeted purpose.
  - Cortex have A, M and R series.

Cortex-A	Cortex-R	Cortex-M
High performance	Real-time response	Low power usage
PCs, laptops, servers, networking equipment, and supercomputers.	medical equipment, vehicle steering, braking and signaling.	IoT and embedded devices, such as wearables and small sensors.



# Micro-architectures of ARM

ISA	Cortex-A	Cortex-M	Cortex-R
ARMv7	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-A7</li> <li>2. Cortex-A9</li> <li>3. Cortex-A17</li> </ol> <b>Products:</b> <ol style="list-style-type: none"> <li>1. Nvidia Tegra 2</li> <li>2. HiSilicon K3V2</li> </ol>	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-M0</li> <li>2. Cortex-M3</li> </ol> <b>Products:</b> <ol style="list-style-type: none"> <li>1. Texas Instruments F28</li> <li>2. Realtek RTL8710</li> </ol>	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-R4</li> <li>2. Cortex-R8</li> </ol>
ARMv8	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-A53</li> <li>2. Cortex-A57</li> <li>3. Cortex-A73</li> </ol> <b>Products:</b> <ol style="list-style-type: none"> <li>1. Apple A7 (iPhone 5)</li> <li>2. Exynos 5433 (Galaxy Note 4)</li> <li>3. Nvidia Tegra X1</li> </ol>	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-M23</li> <li>2. Cortex-M33</li> </ol> <b>Products:</b> <ol style="list-style-type: none"> <li>1. Nuvoton M2351</li> <li>2. GigaDevice GD32E230</li> <li>3. NXP LPC5500</li> <li>4. ST STM32 L5</li> </ol>	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-R52</li> <li>2. Cortex-R82</li> </ol>
ARMv9	<b>Processor cores:</b> <ol style="list-style-type: none"> <li>1. Cortex-A715</li> </ol> <b>Products:</b> <ol style="list-style-type: none"> <li>1. MediaTek Dimensity 9200</li> <li>2. Qualcomm Snapdragon 8 Gen 2</li> </ol>		

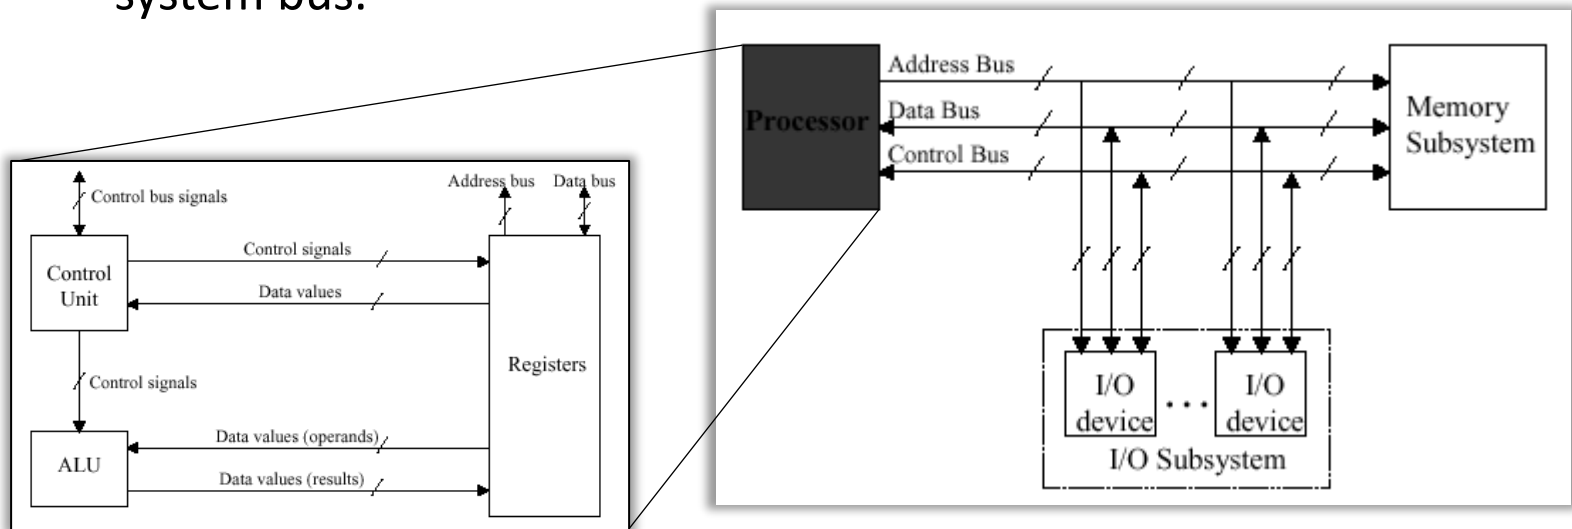
[https://en.wikipedia.org/wiki/Template:Embedded\\_ARM-based\\_chips](https://en.wikipedia.org/wiki/Template:Embedded_ARM-based_chips)

# The Organisation of Microprocessor Systems

Some basic terms and knowledge

# Von Neumann machine

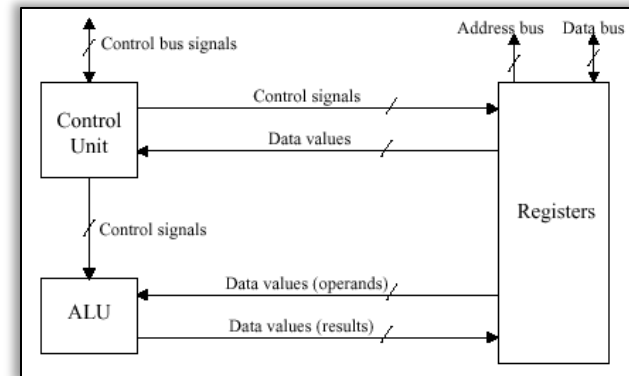
- Below is the typical organization of a modern Von Neumann machine.
  - Data and program are mostly stored in the computer memory separate from the process.
  - Registers in the processor [datapath](#) can also store small amount of data.
  - The address bus, data bus and control bus are also referred to as system bus.







# CPU

- **Registers:** Very fast but limited memory space embedded inside the CPU.
- **Arithmetic logic unit (ALU):** Carries out arithmetic and bitwise operations.
  - Arithmetic: + - \* /
  - Bitwise: & | ^ ~
- **Control unit:** Coordinator of memory, ALU, registers, I/O devices.
- **CPU example workflow:**
  1. Fetch the instruction
  2. Fetch the operands
  3. Execute the instruction
  4. Write the results.
  5. Back to 1.




# Internal Organisation of PC

- The image on the right shows the structure of a PC we see in daily life.
  - You can install operating systems on it.



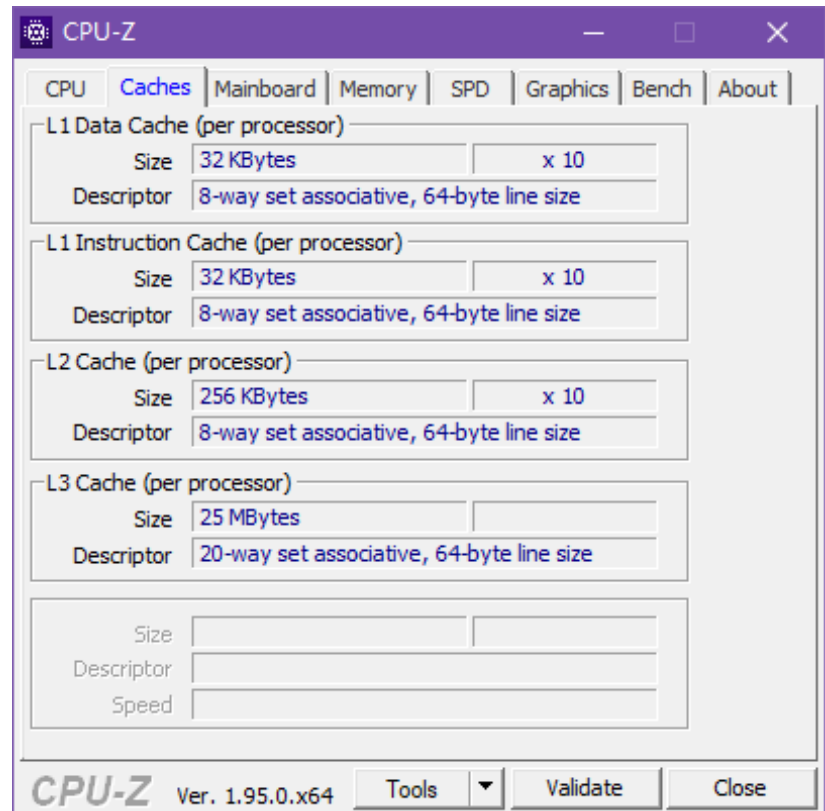
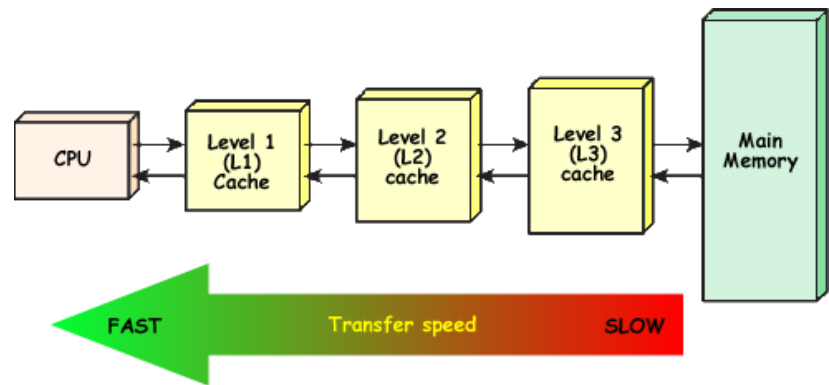
# MCU Board

- The image on the right is an ARM MCU board
  - **MCU**: Microcontroller Unit
  - It only runs one single program.
- You cannot install full operating systems on MCU.
  - Lacks memory management unit (MMU).
  - Some details will be covered later in this module.




# CPU Cache

- Another Component of CPU.
- Lies in between CPU registers and memory.
  - Access speed: Register > cache > memory
  - Manufacture cost: Register > cache > memory
  - Storage capacity: Memory > cache > register
- Before looking for a certain piece of data in memory, CPU will first look for it in the cache.



# Memory



- Data and Program instructions are stored in the **primary memory**.

- A byte is a memory unit for storage
- A memory chip is full of such bytes.
- **Secondary memory** refers to external storages
  - CD, Floppy, Hard disk, USB storage etc.

- Technology limitations affects memory size.

- For 32-bit CPUs

- Address bus is 32-bit wide.
- Each memory location refers to 1 Byte (8 bits)
- Can address up to  $2^{32} = 4294967296 = 4 \text{ GB}$ 
  - The lowest address is  $0 \times 00000000$
  - The highest address is  $0 \times \text{FFFFFFFF}$

1 byte

Address	Content
0xFFFFFFFF	0x8b
0xFFFFFFF	0x11
.	.
.	.
.	.
.	.
0x00000001	0xAA
0x00000000	0xCD

# Accessing Memory

- **Data** and **instructions** are accessed and manipulated in fixed-length chunks.
- **Data sizes in C:** 1 byte (char), 2 bytes (short), 4 bytes (int, float), 8 bytes (long)
  - For data that is less than the size of its `sizeof(type)`, leading zeros/ones will be added.
  - A `char` with value 1 is stored as 000000001 in the memory.
  - A `char` with value -12 is stored as 11110100.
  - A `short` with value -12 is 1111111111110100.
    - Negative binary numbers will be taught later.
- **Data sizes in ARMv7 assembly:** byte, half-word (2 bytes), word (4 bytes)
  - No data type associated.
  - No sign/unsigned number indicator.

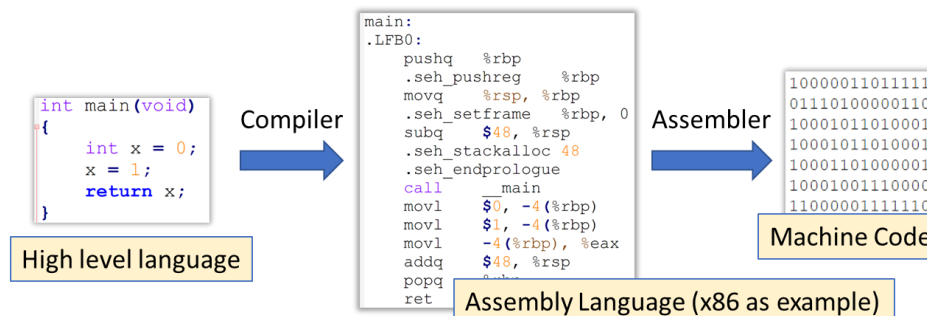
4 bytes

0 ... 0 1

4 x 8 = 32 bit

# Accessing Memory

- Data and instructions are accessed and manipulated in fixed-length chunks.
- Instruction: each instruction (usually) has a length of one **word**.
  - Words in a 32-bit system are usually 32 bits long.
  - Words in a 64-bit system are usually 64 bits long.
    - Somehow related to the width of the data and address bus.
    - BUT, **some ARM instructions are half-word long**.
    - In x86, the length of instructions may vary.
  - Usually, word size = address size (pointer size) = register size.



# Question

- Data in memory has no associated data type information.
- How does a C program know the data type associated with a piece of memory and calculate correct results?

```
int x = 1;
```


```
float y = 1.1;
```

```
x = x * 2;
```

```
y = y * x;
```



# Byte Ordering



- Assume we have a 32-bit int value:

• 0x AA BB CC DD

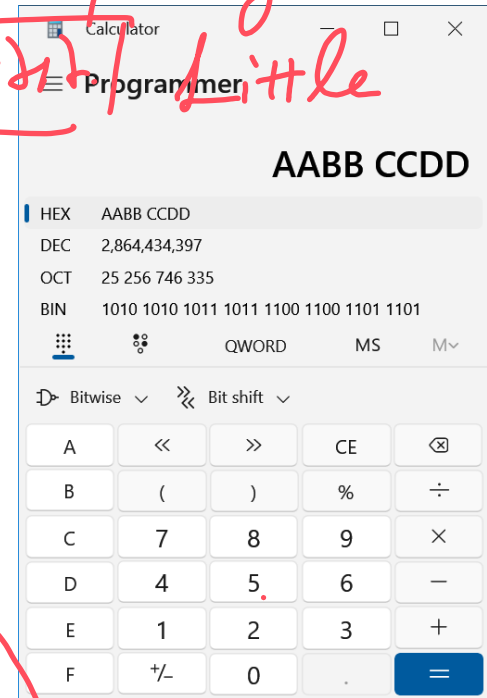
- How should its bytes be ordered in memory?

- Arrangement 1

Value	DD	CC	BB	AA
Address	0x1	0x2	0x3	0x4

- Arrangement 2

Value	AA	BB	CC	DD
Address	0x1	0x2	0x3	0x4



- The ordering of bytes is called **endianness**.
  - The first arrangement method is called **little endian**.
  - The second arrangement method is called **big endian**.

# Byte Ordering

- Big-endian architectures:
  - SPARC, z/Architecture
  - Least significant byte has highest address.
    - DD at 0x04
- Little-endian architectures:
  - x86, x86-64
  - Least significant byte has lowest address.
    - DD at 0x01
- Both endianness are supported:
  - ARM, MIPS, PowerPC
  - Can be specified



# Important Notes about Byte Ordering

- Byte ordering only affects basic data units (word), composite data units like arrays or struct/class are NOT affected by byte ordering.

```
int y[] = {0x41424344, 0x45464748};
```

- Big-endian

Value	0x41	0x42	0x43	0x44	0x45	0x46	0x47	0x48
Address	0x01	0x02	0x03	0x04	0x05	0x06	0x07	0x08

- Little-endian

Value	0x44	0x43	0x42	0x41	0x48	0x47	0x46	0x45
Address	0x01	0x02	0x03	0x04	0x05	0x06	0x07	0x08


# Important Notes about Byte Ordering

“Byte ordering only affects basic data units (word), composite data units like arrays or struct/class are NOT affected by byte ordering.”

- How would you prove this?
  - Assume we use x86, it should be little-endian.
  - Write a piece of C/C++ code

# Alignment

- Instructions that read/write data from/to memory views the memory as a series of chunks of fixed-sized data.
- Each chunk has:
  - The address (of the first byte) of that chunk of data.
  - The length of data.
- If an instruction loads 4 bytes in one go, then the address of the 4-byte data must be a multiply of 4.
- This is called memory **alignment**.
  - Aligned vs unaligned access.



# Alignment: A Case Study

- How much memory does struct t1 occupy?

```
struct t1 {  
    int a;  
    char b1;  
    char b2;  
    char b3;  
    char b4;  
    int c;  
};
```

```
struct t2 {  
    int a;  
    char b1;  
    char b2;  
    int c;  
};
```

- How about struct t2?

# Alignment: A Case Study

- If t1 and t2 are initialised as follows:

```
struct t1 a = {1, 2, 3, 4, 5, 6};  
struct t2 b = {1, 2, 3, 4};
```

- How is the arrangement of these numbers in the memory?
- Try the code below for 'a':

```
char * c = (char *) &a;  
printf("showing struct t1 a\n");  
for (int i = 0; i < 12; i++) {  
    printf("byte %d: %d\n", i, *c);  
    c++;  
}
```

# Alignment: A Case Study

- The second chunk of 32 bits of memory in struct t2 only stores 2 characters, the rest of it is left empty.
- If bytes 6~9 are used to store int, the access will be unaligned

showing struct t1 a	showing struct t2 b
byte 0: 1	byte 0: 1
byte 1: 0	byte 1: 0
byte 2: 0	byte 2: 0
byte 3: 0	byte 3: 0
byte 4: 2	byte 4: 2
byte 5: 3	byte 5: 3
byte 6: 4	byte 6: 0
byte 7: 5	byte 7: 0
byte 8: 6	byte 8: 4
byte 9: 0	byte 9: 0
byte 10: 0	byte 10: 0
byte 11: 0	byte 11: 0



# Extended Exploration

Which of the following questions pique your curiosity?

- What's the difference between the Von-Neumann architecture and the Harvard architecture?
- How do CPUs carry out instructions so fast?
- How are function calls in C/C++ translated into assembly?
- Do I waste memory by declaring local variables in functions but then never use them?
- Why I sometimes see values like `3.1500000000000000004` but it should really be `3.15` when using `printf()`?
- What happens when I assign an object to another object in C++?
- How I get virus after visiting some malicious website?